

M146 Database Systems Spring 2021

Georgia Koutrika



Transactions in SQL

- In “ad-hoc” SQL:
 - Default: each statement = one transaction
- In a program, multiple statements can be grouped together as a transaction:

```
START TRANSACTION
    UPDATE Bank SET amount = amount - 3000
    WHERE name = 'Bob'
    UPDATE Bank SET amount = amount + 3000
    WHERE name = 'Joe'
COMMIT
```

Motivation for Transactions

Grouping user actions (reads & writes) into *transactions* helps with two goals:

1. **Recovery & Durability**: Keeping the DBMS data consistent and durable in the face of crashes, aborts, system shutdowns, etc.
2. **Concurrency**: Achieving better performance by parallelizing TXNs *without* creating anomalies

Motivation for Transactions

1. Recovery & Durability of user data is essential for reliable DBMS usage

- The DBMS may experience crashes (e.g. power outages, etc.)
- Individual TXNs may be aborted (e.g. by the user)

Idea:

Make sure that TXNs are either **durably stored in full, or not at all**;
Keep log to be able to “roll-back” TXNs

Motivation for Transactions

2. Concurrent execution of user programs is essential for good DBMS performance.

- Disk accesses may be frequent and **slow**
optimize for throughput (# of TXNs), trade for latency (time for any one TXN)
- Users should still be able to execute TXNs as if in **isolation** and such that **consistency** is maintained

Idea: Have the DBMS handle running several user TXNs concurrently, in order to keep CPUs busy...
The DBMS uses locks to ensure correctness

Transaction Properties: ACID

- **Atomicity**: all the actions in a transaction are executed as a single atomic operation; either they are all carried out or none are
- **Consistency**: if a transaction begins with the DB in a consistent state, it must finish with the DB in a consistent state
- **Isolation**: a transaction should execute as if it is the only one executing; it is protected (isolated) from the effects of concurrently running transactions
- **Durability**: if a transaction has been successfully completed, its effects should be permanent

ACID continues to be a source of great debate!

Ensuring Atomicity & Durability

ACID

- Atomicity:
 - TXNs should either happen completely or not at all
 - If abort / crash during TXN, *no* effects should be seen
- Durability:
 - If DBMS stops running, changes due to completed TXNs should all persist
 - *Just store on stable disk*

TXN 1

Crash / abort

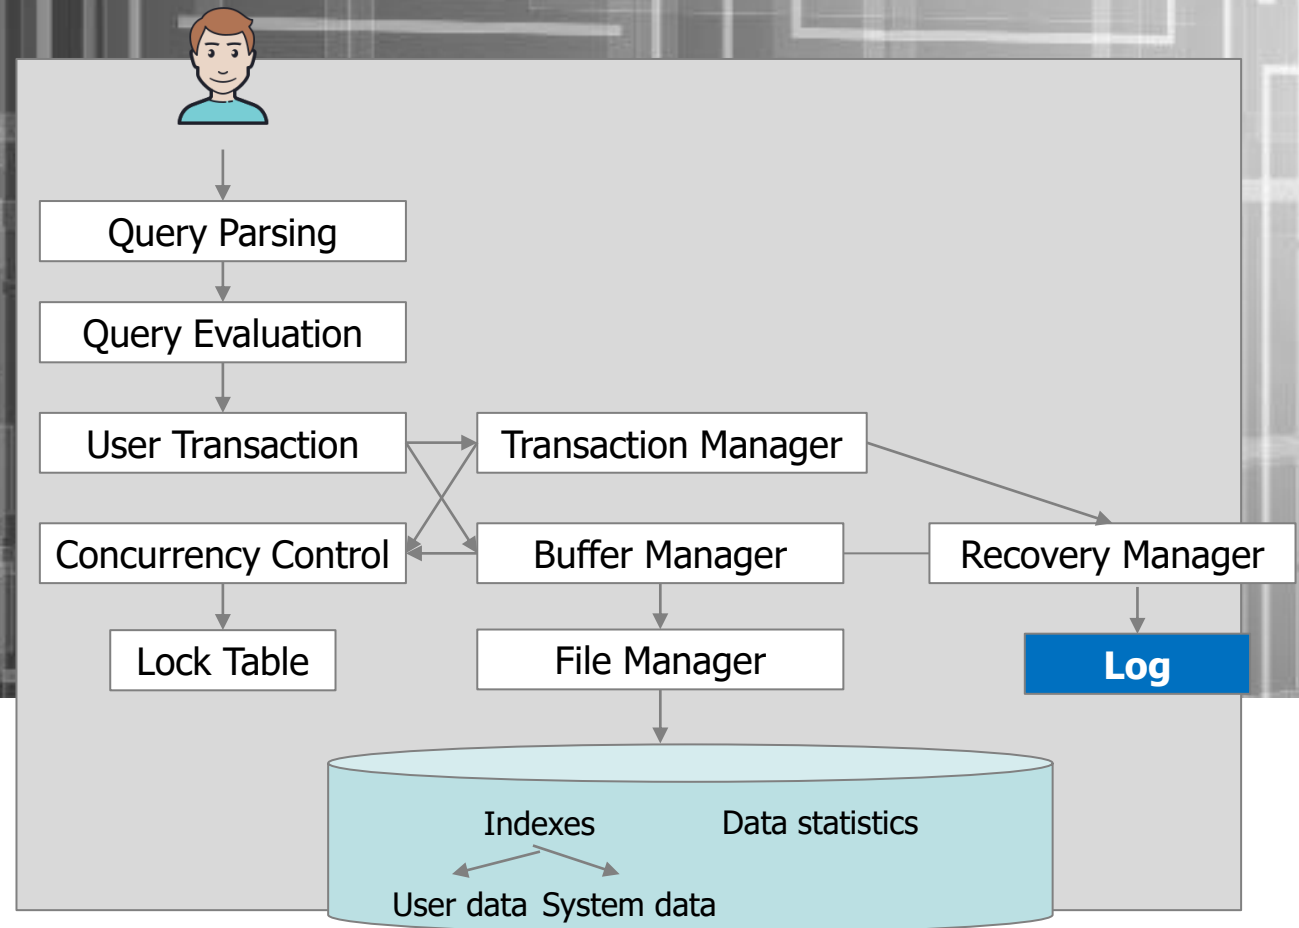
No changes
persisted

TXN 2

All changes
persisted

We'll focus on how to accomplish atomicity (via logging)

LOGGING



The log

- The following actions are recorded in the log
 - Whenever a transaction writes an object
 - Whenever a transaction commits/aborts
- The log record must be on disk before the data record reaches the disk

The Log

- Log is *duplexed* and *archived* on stable storage.
- Can **force write** entries to disk
 - A page goes to disk.
- All log-related activities (in fact, all concurrency control related activities) are handled by the DBMS
- The user does not know anything

Basic Idea: (Physical) Logging

- Record UNDO information for every update!
 - Sequential writes to log
 - Minimal info (diff) written to log
- The **log** consists of an **ordered list of actions**
 - Log record contains:
 <XID, location, old data, new data>
 Log records are chained by transaction ID (why?)

This is sufficient to UNDO any transaction!


Why do we need logging for atomicity?

- Couldn't we just write TXN to disk **only** once whole TXN complete?
 - Then, if abort / crash and TXN not complete, it has no effect- atomicity!
 - *With unlimited memory and time, this could work...*
- However, we **need to log partial results of TXNs** because of:
 - Memory constraints (space for full TXN??)
 - Time constraints (what if one TXN takes very long?)

We need to write partial results to disk!
...And so we need a **log** to be able to ***undo*** these partial results!

A picture of logging

T: R(A), W(A)



A picture of logging

T: R(A), W(A)

A: 0 → 1



Operation recorded in
log in main memory!

What is the correct
way to write this all to
disk?



What is the correct way to write this all to disk?

- We'll look at the *Write-Ahead Logging (WAL)* protocol
- We'll see why it works by looking at other protocols which are incorrect!

Remember: Key idea is to ensure durability *while* maintaining our ability to “undo”!

Transaction Commit Process

1. FORCE Write **commit** record to log
2. All log records up to last update from this TX are FORCED
3. Commit() returns

Transaction is committed *once commit log record is on stable storage*

Incorrect Commit Protocol #1

T: R(A), W(A)

A: 0 → 1



Let's try committing *before* we've written either data or log to disk...

OK, Commit!

If we crash now, is T durable?



Lost T's update!

Incorrect Commit Protocol #2

T: R(A), W(A)

A: 0 → 1



Let's try committing *after* we've written data but *before* we've written log to disk...

OK, Commit!

If we crash now, is T durable? Yes! Except...



How do we know whether T was committed??

Write-ahead Logging (WAL) Commit Protocol

T: R(A), W(A)

A: 0 → 1



This time, let's try committing **after we've written log to disk but before we've written data to disk**... this is WAL!

OK, Commit!

If we crash now, is T durable?



Write-ahead Logging (WAL) Commit Protocol

T: R(A), W(A)



This time, let's try committing **after we've written log to disk but before we've written data to disk**... this is WAL!

OK, Commit!

If we crash now, is T durable?



A: 0 → 1



USE THE LOG!

Write-Ahead Logging (WAL)

- DB uses **Write-Ahead Logging (WAL)** Protocol:

Each update is logged! Why not reads?

1. Must *force log record* for an update *before* the corresponding data page goes to storage

→ **Atomicity**

2. Must *write all log records* for a TX *before commit*

→ **Durability**

Logging Summary

- If DB says TX **commits**, TX effect **remains** after database crash
- DB can **undo actions** and help us with **atomicity**
- This is only half the story...

Crash recovery

Three phases to recovery (ARIES)

- **Analysis**: scan log forward, identifying committed and aborted/unfinished transactions
- **Redo**: all committed transactions are made durable
- **Undo**: the actions of all aborted and/or unfinished transactions are undone