

M146 Database Systems Spring 2021

Georgia Koutrika



Query Optimizer

- What it needs:
 1. Information about how to compute the relational operators in the tree
 - Based on the access paths and algorithms available
 2. Information about the data stored
 - System Catalog Information
 3. Formulas to compute cardinalities and costs
 4. Strategy to generate plans and select the one to be executed

System Catalog Information

– Information about the size of a file

- n_R : number of tuples in a relation R .
- b_R : number of blocks containing tuples of R .
- l_R : record size of R .
- bf_R : blocking factor of R — i.e., the number of tuples of R that fit into one block.

– Information about indexes and indexing attributes of a file

- Number of levels (x) of each multilevel index
- Number of first-level index blocks (b_{i1})
- Number of distinct values (d) of an attribute
- Selectivity (sl) of an attribute

Example (System Catalog)

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.

In reality, the system catalog ...

- ACCESS_POLICY
- ALL_TABLES
- AUDIT_MANAGING_USERS_PRIVILEGES
- CATALOG_SUBSCRIPTION_CHANGES
- CATALOG_SYNC_STATE
- CATALOG_TRUNCATION_STATUS
- CLIENT_AUTH
- CLIENT_AUTH_PARAMS
- CLUSTER_LAYOUT
- COLUMNS
- COMMENTS
- CONSTRAINT_COLUMNS
- DATABASES
- DIRECTED_QUERIES
- DUAL
- ELASTIC_CLUSTER
- EPOCHS
- FAULT_GROUPS
- FOREIGN_KEYS
- GRANTS
- HCATALOG_COLUMNS
- HCATALOG_SCHEMATA
- HCATALOG_TABLES
- HCATALOG_TABLE_LIST
- KEYWORDS
- LARGE_CLUSTER_CONFIGURATION_STATUS
- LICENSE_AUDITS
- LICENSES

- LOAD_BALANCE_GROUPS
- LOG_PARAMS
- LOG_QUERIES
- LOG_TABLES
- MATERIALIZE_FLEXTABLE_COLUMNS_RESULTS
- MODELS
- NETWORK_ADDRESSES
- NODES
- NODE_SUBSCRIPTION_CHANGE_PHASES
- NODE_SUBSCRIPTIONS
- ODBC_COLUMNS
- PASSWORD_AUDITOR
- PASSWORDS
- PRIMARY_KEYS
- PROFILE_PARAMETERS
- PROFILES
- PROJECTION_CHECKPOINT_EPOCHS
- PROJECTION_COLUMNS
- PROJECTION_DELETE_CONCERNS
- PROJECTIONS
- RESOURCE_POOL_DEFAULTS
- RESOURCE_POOLS
- ROLES
- ROUTING_RULES
- SCHEMATA
- SEQUENCES
- SESSION_SUBSCRIPTIONS
- SHARDS
- STORAGE_LOCATIONS
- SYSTEM_COLUMNS
- SYSTEM_TABLES
- TABLE_CONSTRAINTS
- TABLES

-
- PROJECTIONS
 - RESOURCE_POOL_DEFAULTS
 - RESOURCE_POOLS
 - ROLES
 - ROUTING_RULES
 - SCHEMATA
 - SEQUENCES
 - SESSION_SUBSCRIPTIONS
 - SHARDS
 - STORAGE_LOCATIONS
 - SYSTEM_COLUMNS
 - SYSTEM_TABLES
 - TABLE_CONSTRAINTS
 - TABLES
 - TEXT_INDICES
 - TYPES
 - USER_AUDITS
 - USER_CLIENT_AUTH
 - USER_FUNCTION_PARAMETERS
 - USER_FUNCTIONS
 - USER_PROCEDURES
 - USER_TRANSFORMS
 - USERS
 - VIEW_COLUMNS
 - VIEW_TABLES
 - VIEWS

Vertica

How to compute the cost of an execution plan?

Let us now see a simple example. We will need:

- statistics
- formulas for the costs of operations
- formulas for cardinalities

Example

Let us assume that we have these tables:

PROJECT(Pnumber, Plocation, Dnum, PStartDate)

DEPARTMENT(Dnumber, Dname, Mgr_ssn)

EMPLOYEE(SSN, Fname, Lname, Address, Bdate)

SQL:

SELECT P.Pnumber, P.Dnum, E.Lname, E.Address, E.Bdate

FROM PROJECT P, DEPARTMENT D, EMPLOYEE E

WHERE P.Dnum = D.Dnumber AND D.Mgr_ssn = E.SSN

AND P.Plocation = 'STAFFORD';

Example

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

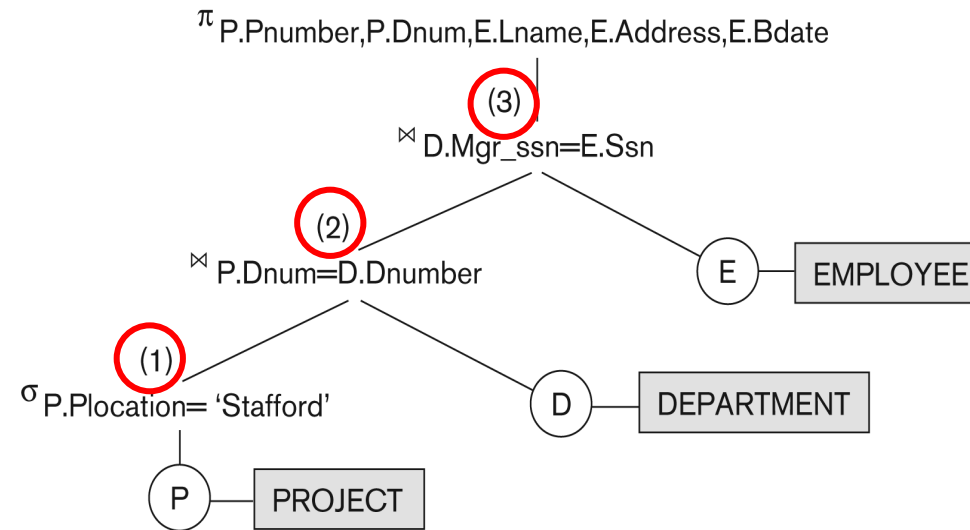
(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.



Example

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

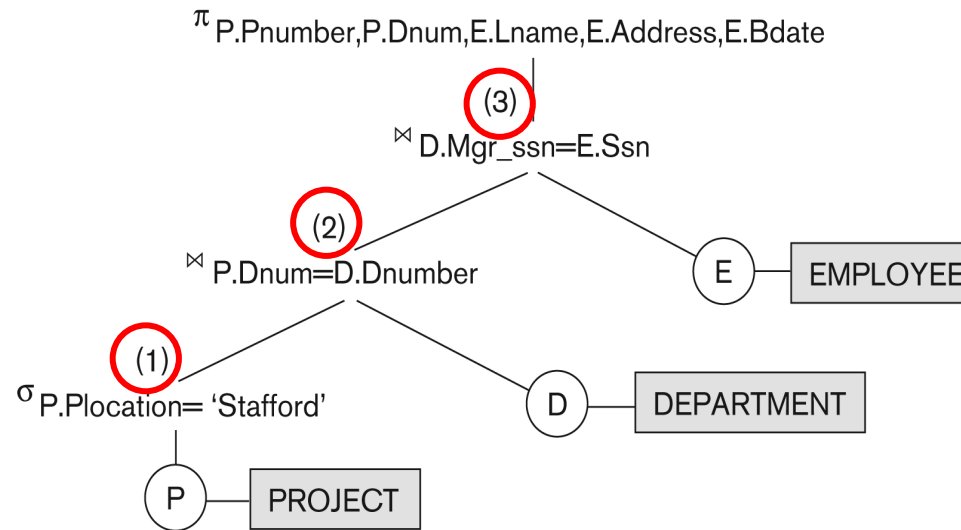
(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.



1 $\sigma_{Plocation = 'Stafford'}$ (PROJECT)

- Table scan (Plocation is not primary key)
 - Cost = 100
- PROJ_PLOC Index (number of levels, $x = 2$)
 - Selectivity = $1/200$ (assuming uniformly distributed)
 - Selection cardinality = Selectivity * Num_rows = 10 rows \rightarrow 10 blocks
 - Cost = $2 + 10 = 12$ ✓

Example

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

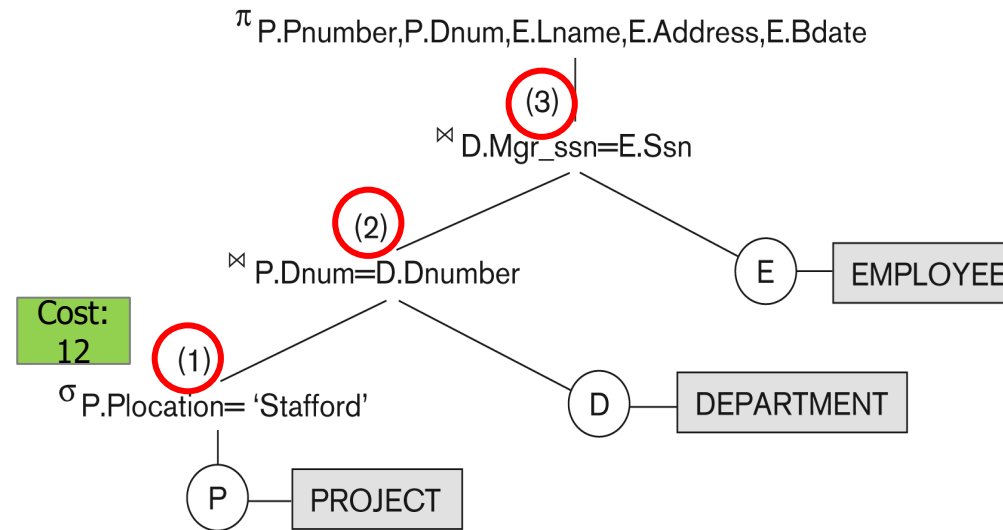
(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.



1 $\sigma_{Plocation = 'Stafford'}$ (PROJECT)

- Table scan (Plocation is not primary key)
 - Cost = 100
- PROJ_PLOC Index (number of levels, x = 2)
 - Selectivity = 1/200 (assuming uniformly distributed)
 - Selection cardinality = Selectivity * Num_rows = 10 rows → 10 blocks
 - Cost = 2 + 10 = 12 ✓

Example

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

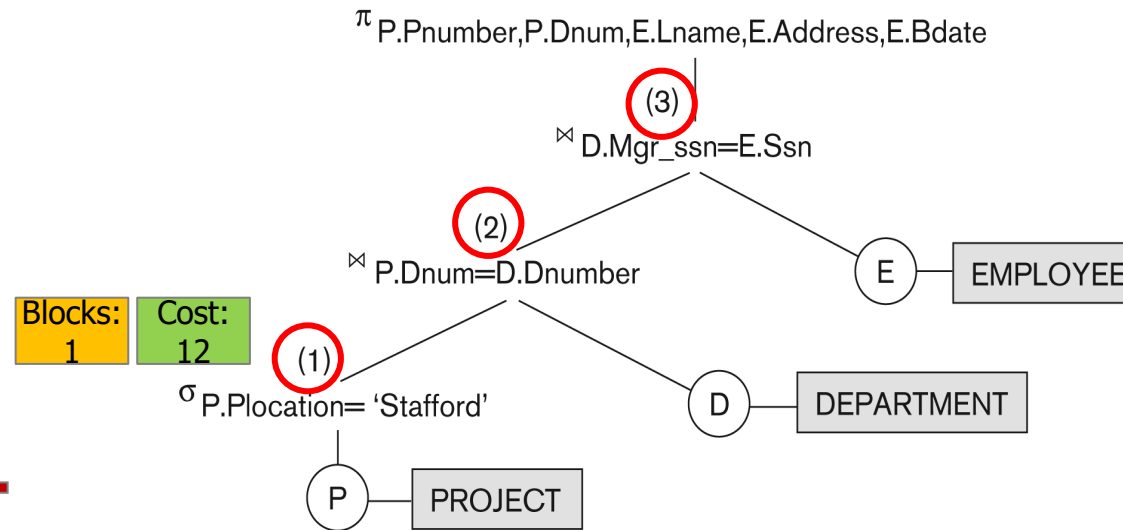
(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.



1 $\sigma_{Plocation = 'Stafford'} (PROJECT)$ = TEMP1

- Estimated number of rows = $2000/200 = 10$
- Blocking factor = $2000/100 = 20$ tuples/block
- So, number of blocks needed = 1

Example

(a)

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |
| EMPLOYEE | Ssn | 10000 | 1 | 10000 |
| EMPLOYEE | Dno | 50 | 1 | 50 |
| EMPLOYEE | Salary | 500 | 1 | 500 |

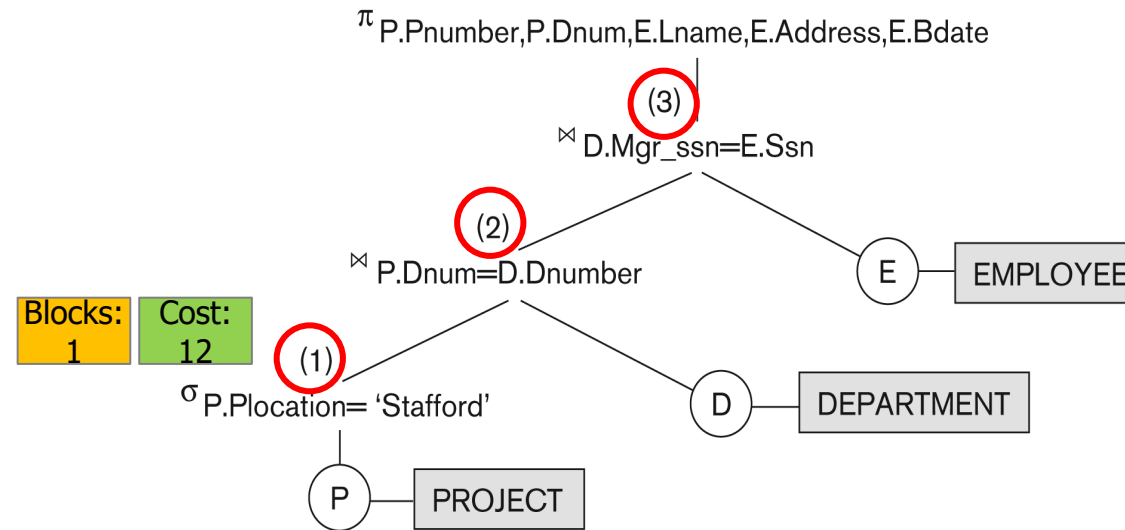
(b)

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

(c)

| Index_name | Uniqueness | Blevel* | Leaf_blocks | Distinct_keys |
|------------|------------|---------|-------------|---------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 | 200 |
| EMP_SSN | UNIQUE | 1 | 50 | 10000 |
| EMP_SAL | NONUNIQUE | 1 | 50 | 500 |

*Blevel is the number of levels without the leaf level.



TEMP1

2 Cost for $\sigma_{Plocation = 'Stafford'} (PROJECT) \bowtie DEPARTMENT$

- No index available to process the join
- We use the nested loop join

$$P(R) + T(R) * P(S) + OUT$$

Example

| Table_name | Column_name | Num_distinct | Low_value | High_value |
|------------|-------------|--------------|-----------|------------|
| PROJECT | Plocation | 200 | 1 | 200 |
| PROJECT | Pnumber | 2000 | 1 | 2000 |
| PROJECT | Dnum | 50 | 1 | 50 |
| DEPARTMENT | Dnumber | 50 | 1 | 50 |
| DEPARTMENT | Mgr_ssn | 50 | 1 | 50 |

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

2 Nested loop join TEMP1 $\bowtie_{Dnum=Dnumber}$ DEPARTMENT = TEMP2

– TEMP1: result of $\sigma_{Plocation = 'Stafford'}(PROJECT)$

- Estimated number of rows = $2000/200 = 10$
- Number of blocks needed = 1

– DEPARTMENT

- number of blocks needed = 5

$P(R) + T(R) * P(S) + OUT$

Example

2 Nested loop join TEMP1 $\bowtie_{Dnum=Dnumber}$ DEPARTMENT = TEMP2

Cost:
53

$$P(R) + T(R)*P(S) + OUT$$

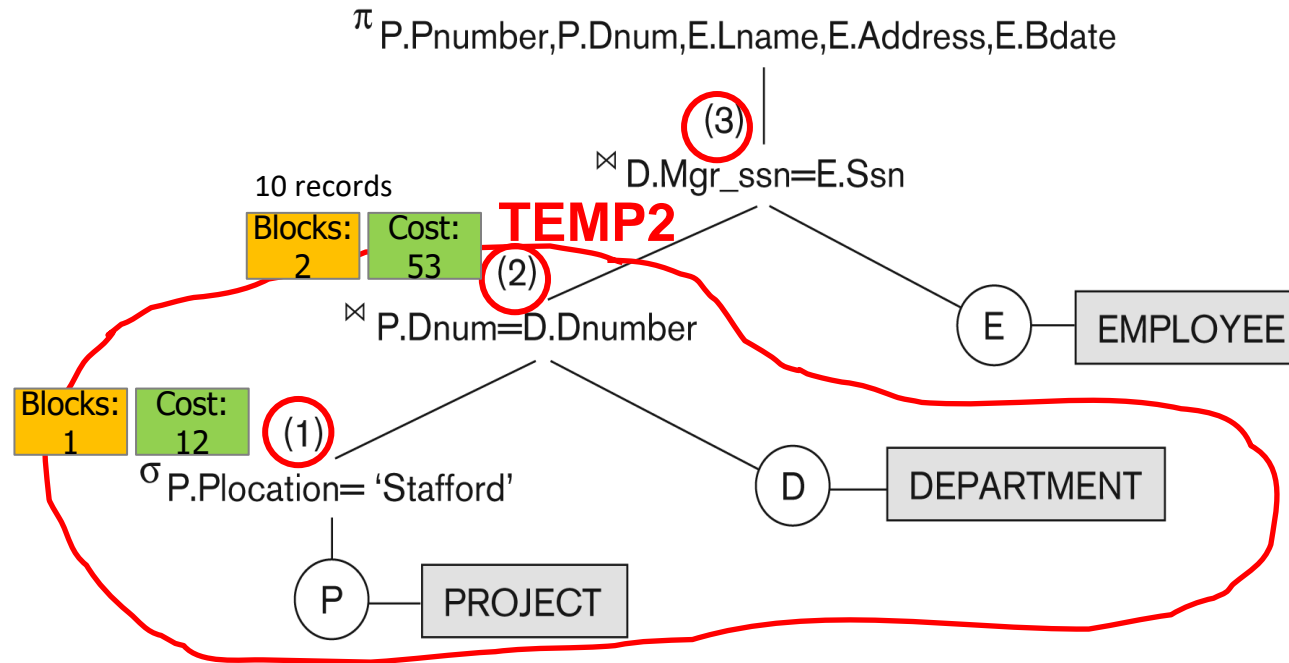
– Use TEMP1 in outer loop for nested-loop join

- Cost = $1 + 10*5$ + cost to write join output into TEMP2
= 51 + cost to write join output into TEMP2

– What is the cost for writing join output?

- Each row in TEMP1 joins exactly with 1 row in DEPARTMENT (why?)
- Estimated number of rows in TEMP2 = 10
join attribute Dnumber is the key of department.
So we assume there are 10 joined records
- Estimated blocking factor = 5 (from estimated record size)
- Number of blocks needed = 2

Example



③ Cost for join TEMP2 \bowtie Mgr_ssn=Ssn EMPLOYEE

Example

| Table_name | Column_name | Num_distinct |
|------------|-------------|--------------|
| PROJECT | Plocation | 200 |
| PROJECT | Pnumber | 2000 |
| PROJECT | Dnum | 50 |
| DEPARTMENT | Dnumber | 50 |
| DEPARTMENT | Mgr_ssn | 50 |
| EMPLOYEE | Ssn | 10000 |
| EMPLOYEE | Dno | 50 |
| EMPLOYEE | Salary | 500 |

| Table_name | Num_rows | Blocks |
|------------|----------|--------|
| PROJECT | 2000 | 100 |
| DEPARTMENT | 50 | 5 |
| EMPLOYEE | 10000 | 2000 |

| Index_name | Uniqueness | Blevel* | Leaf_blocks |
|------------|------------|---------|-------------|
| PROJ_PLOC | NONUNIQUE | 1 | 4 |
| EMP_SSN | UNIQUE | 1 | 50 |
| EMP_SAL | NONUNIQUE | 1 | 50 |

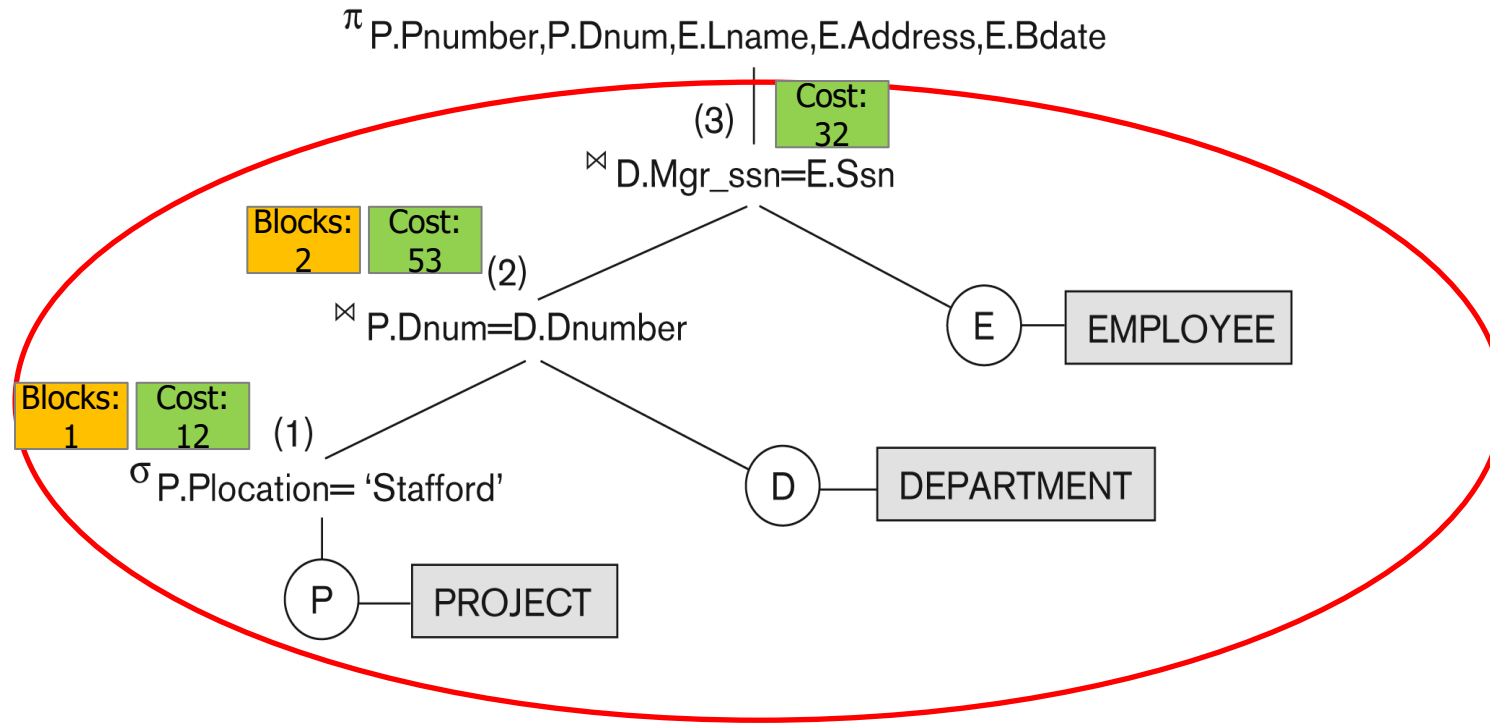
*Blevel is the number of levels without the leaf level.

③ Cost for join TEMP2 $\bowtie_{Mgr_ssn=Ssn}$ EMPLOYEE

$$P(R) + T(R) * L + OUT$$

- Primary index (EMP_SSN) available for Ssn in EMPLOYEE
- Can use Index Nested Loop join on TEMP2
- For each row in TEMP2, use primary index to retrieve corresponding rows in EMPLOYEE
 - Cost = 2 + 10 × (1 + 1 + 1) + cost of output = 32 + cost of output

Example



- Use pipelining to produce the final result
 - So, no additional cost for projection
 - Total cost = 12 + 53 + 32 + cost of writing final output

Query Optimizer

- What it needs:
 1. Information about how to compute the relational operators in the tree
 - Based on the access paths and algorithms available ✓
 2. Information about the data stored ✓
 - System Catalog Information
 3. **Formulas to compute costs and cardinalities**
 4. Strategy to generate plans and select the one to be executed

Selectivity

Consider a query predicate, such as `WHERE last_name LIKE 'A%'` (or a combination of predicates)

Selectivity is the **percentage of rows** returned by a query predicate

- with 0 meaning no rows
- 1 meaning all rows.

A predicate becomes more selective as the selectivity value approaches 0 and less selective (or more unselective) as the value approaches 1.

Selectivity

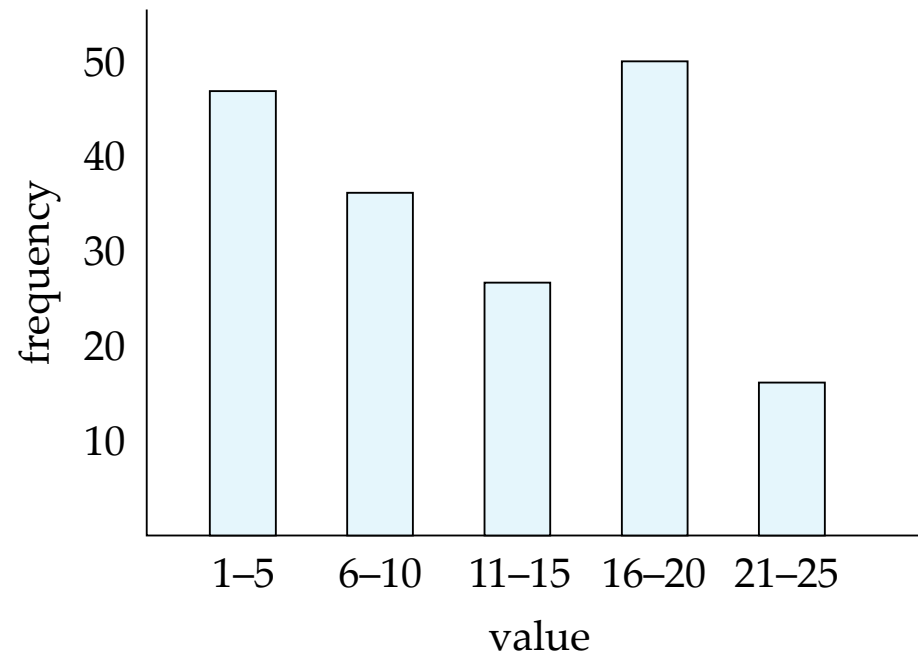
For an equality predicate :

Selectivity = $1/(\text{number of distinct values})$

If there is a **histogram** on a column, then the estimator uses the histogram instead of the number of distinct values. The histogram captures the distribution of different values in a column, so it yields better selectivity estimates, especially for columns that have data skew.

Histograms

- Histogram on attribute *age* of relation *person*



- **Equi-width** histograms
- **Equi-depth** histograms

Cardinality

The **cardinality** is the **number of rows** returned by each operation in an execution plan. This input, which is **crucial to obtaining an optimal plan**, is common to all cost functions.

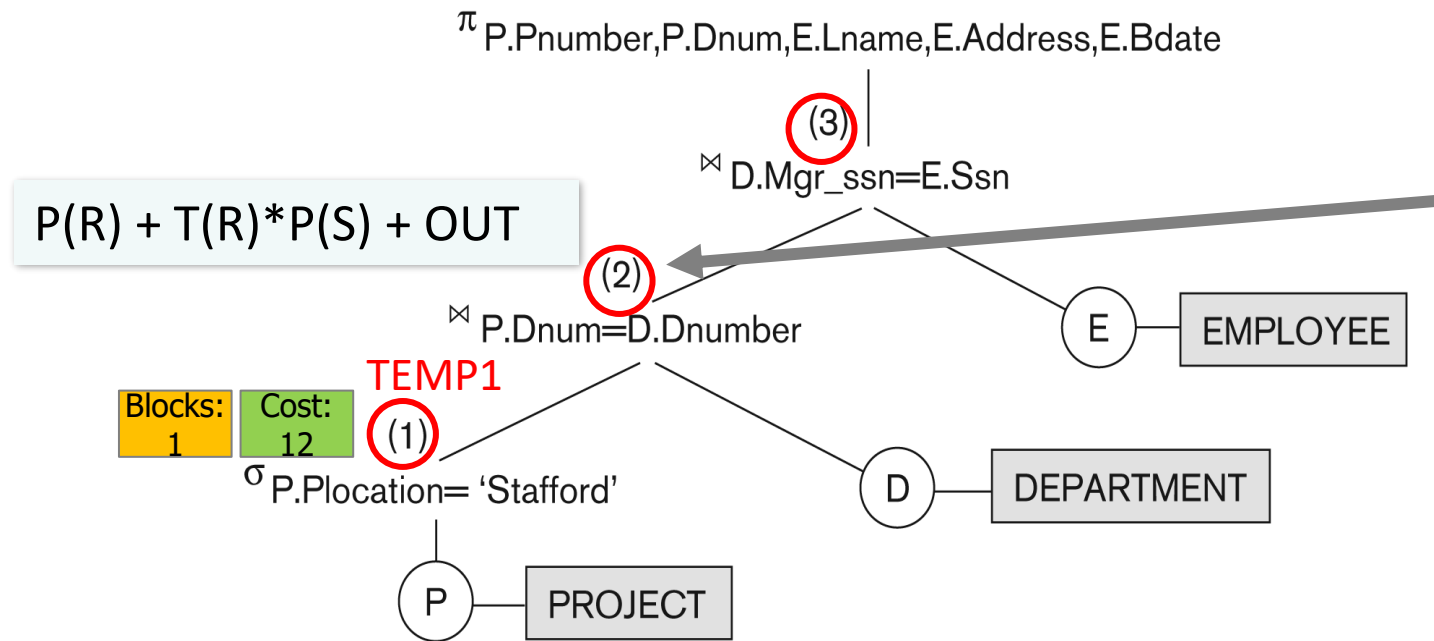
The estimator can derive cardinality from:

- the table statistics,
- after accounting for effects from predicates (filter, join, and so on), DISTINCT or GROUP BY operations, and so on.

Cardinality for a SQL query with 1 equality predicate = (number of rows)/(number of distinct values)

Cardinality

Cardinality estimates must be **as accurate as possible** because they influence all aspects of the execution plan. Cardinality is important when the optimizer determines the cost of a join, the cost of a sort.



For example:

For operation 2 (the join),
we picked TEMP1 as the outer relation!

Cardinality Estimation of Query Result

There are two principal approaches to query cardinality estimation:

Database Profile.

- Maintain **statistical information** about numbers and sizes of tuples, distribution of attribute values for base relations, as part of the database catalog (meta information) during database updates.
- Calculate these parameters for **intermediate query results** based upon a (simple) statistical model during query optimization.
- Typically, the statistical model is based upon the **uniformity** and **independence** assumptions.
- Both are typically not valid, but they allow for simple calculations → limited accuracy.
- In order to improve accuracy, the system can record **histograms** to more closely model the actual value distributions in relations.

Cardinality Estimation of Query Result

There are two principal approaches to query cardinality estimation:

Sampling Techniques.

- Gather the necessary characteristics of a query plan (base relations and intermediate results) at **query execution time**
 - **Run query on a small sample** of the input.
 - **Extrapolate** to the full input size.
 - It is crucial to find the **right balance** between sample size and the resulting accuracy

Statistical Information

- n_R : number of tuples in a relation R .
- b_R : number of blocks containing tuples of R .
- l_R : record size of R .
- bf_R : blocking factor of R — i.e., the number of tuples of R that fit into one block.
- $V(A, R)$: number of distinct values that appear in R for attribute A
- If tuples of R are stored together physically in a file, then:

$$b_R = n_R / bf_R$$

Assumptions

In order to obtain tractable cardinality estimation formulae, assume one of the following:

Uniformity & independence (simple, yet rarely realistic)

All values of an attribute **uniformly appear** with the same probability.
Values of different attributes **are independent** of each other.

Worst case (unrealistic)

No knowledge about relation contents at all.

Perfect knowledge (unrealistic)

Details about the exact distribution of values are known.
Requires huge catalog or prior knowledge of incoming queries.

Cardinality Estimation of Selection

- $\sigma_{A=v}(R)$
 - **cardinality** = $n_R / V(A,R)$: number of records that will satisfy the selection
 - Equality condition on a key attribute: **cardinality** = 1

$V(A, R)$: number of distinct values that appear in R for attribute A

Uniformity

Cardinality Estimation of Selection

- $\sigma_{A=v}(R)$
 - **cardinality** = $n_R / V(A,R)$: number of records that will satisfy the selection
 - Equality condition on a key attribute: **cardinality** = 1
- $\sigma_{A \leq v}(R)$ (case of $\sigma_{A \geq v}(R)$ is symmetric)
 - If $\min(A,R)$ and $\max(A,R)$ are available in catalog
 - **cardinality** = 0 if $v < \min(A,R)$
 - **cardinality** =
$$n_R \cdot \frac{v - \min(A,R)}{\max(A,R) - \min(A,R)}$$
 - If histograms available, we can refine above estimate
 - In absence of statistical information **cardinality** is assumed to be $n_R / 2$.

Uniformity

Cardinality of Complex Selections

The **selectivity** of a condition θ_i is the **probability** that a tuple in the relation r satisfies θ_i .

- If s_i is the number of satisfying tuples in r ,
selectivity of $\theta_i = s_i / n_r$

Uniformity & independence

Let us recall some probability formulas:

$$P(A \text{ and } B) = P(A) \cdot P(B)$$

Conjunction: $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)$. Assuming independence,

$$\text{cardinality} = n_r \cdot \frac{s_1 \cdot s_2 \cdot \dots \cdot s_n}{n_r^n}$$

Cardinality of Complex Selections

The **selectivity** of a condition θ_i is the **probability** that a tuple in the relation r satisfies θ_i .

- If s_i is the number of satisfying tuples in r ,
selectivity of $\theta_i = s_i / n_r$

Uniformity & independence

Let us recall some probability formulas:

$$P(A \text{ OR } B) = P(A) + P(B) - P(A) \cdot P(B)$$

Disjunction: $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$.

$$\text{cardinality} = n_r \cdot \left(1 - \left(1 - \frac{s_1}{n_r} \right) \cdot \left(1 - \frac{s_2}{n_r} \right) \cdot \dots \cdot \left(1 - \frac{s_n}{n_r} \right) \right)$$

Negation: $\sigma_{\neg \theta}(r)$.

$$\text{cardinality} = n_r - \text{size}(\sigma_{\theta}(r))$$

Cardinality of Projections

$\Pi_A(R)$

cardinality = $\left\{ \begin{array}{ll} V(A, R) & \text{if } A \text{ is a single attribute} \\ |R| & \text{if } A \text{ contains the key of } R \\ \text{Min}(|R|, \Pi_{A_i} V(A_i, R)) & \text{otherwise} \end{array} \right.$

Independence

$V(A, R)$: number of distinct values that appear in R for attribute A

Cardinality of Joins

- If $R \cap S = \{A\}$ is not a key for R or S .

If we assume that **every tuple t in R produces tuples in $R \bowtie S$** ,
the number of tuples in $R \bowtie S$ is estimated to be:

$$\text{cardinality} = \frac{n_r * n_s}{V(A, s)}$$

If the reverse is true, the estimate obtained will be:

$$\text{cardinality} = \frac{n_r * n_s}{V(A, r)}$$

The lower of these two estimates is probably the more accurate one.

- Can improve on above if histograms are available
 - Use formula similar to above, for each cell of histograms on the two relations

n_R : number of tuples in a relation R .

$V(A, R)$: number of distinct values that appear in R for attribute A

Cardinality of Other Operations

- Aggregation : **cardinality** of $_{Ag_F}(r)$ = $V(A,r)$

Cardinality of Other Operations

- Set operations

- For unions/intersections of selections on the same relation:
rewrite and use size estimate for selections

- E.g. $\sigma_{\theta_1}(r) \cup \sigma_{\theta_2}(r)$ can be rewritten as $\sigma_{\theta_1 \vee \theta_2}(r)$

-

For operations on different relations:

- **cardinality** of $r \cup s$ = size of r + size of s .
- **cardinality** of $r \cap s$ = minimum size of r and size of s .
- **cardinality** of $r - s$ = r .

All the three estimates may be quite inaccurate, but provide upper bounds on the sizes.

Query Optimizer

- What it needs:
 1. Information about how to compute the relational operators in the tree
 - Based on the access paths and algorithms available
 2. Information about the data stored
 - System Catalog Information
 3. Formulas to compute costs and cardinalities
 4. **Strategy to generate plans and select the one to be executed**

Plan Generator

The plan generator explores various plans for a query block by trying out **different access paths, join methods, and join orders**.

Many plans are possible because of the various combinations that the database can use to produce the same result.

The optimizer picks the plan with the **lowest** cost (from the ones it examines)

Search Space Challenges

For example: Join orders

Consider **finding the best join-order** for $r_1 \ r_2 \ \dots \ r_n$.

There are $(2(n-1))!/(n-1)!$ different join orders for above expression.

Search space is huge!

- Many possible equivalent trees
- Many implementations for each operator
- Many access paths for each relation
- Cannot consider ALL plans
- Want a search space that includes low-cost plans

A local optimal method

- Choose the best algorithm for each operator
- The global effect may not be optimal
- Must consider the **interaction of evaluation techniques** when choosing evaluation plans
 - choosing the cheapest algorithm for each operation independently may not yield best overall cost. E.g.
 - merge-join may be costlier than hash-join, but may provide a sorted output which reduces the cost for an outer level aggregation.
 - nested-loop join may provide opportunity for pipelining

Choice of Evaluation Plans

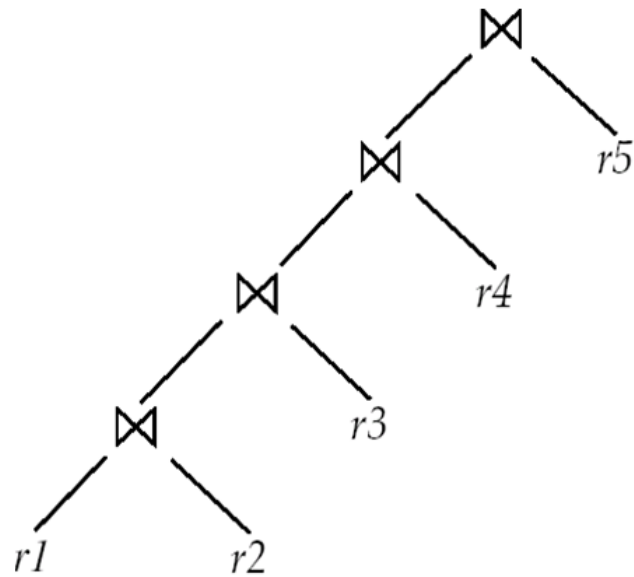
- Practical query optimizers incorporate elements of the following two broad approaches:
 1. Search all the plans and choose the best plan in a cost-based fashion.
 2. Uses heuristics to choose a plan.

System R Search Space

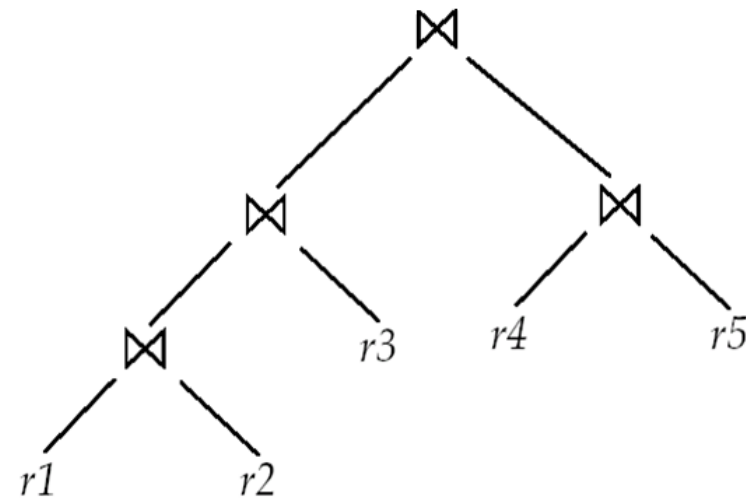
- **Only left-deep plans** ①
 - Enable **dynamic programming** ② for enumeration
 - Facilitate **tuple pipelining** from outer relation
- Consider plans with all **“interesting orders”** ③
- Perform cross-products after all other joins (heuristic)
- Only consider nested loop & sort-merge joins
- Consider both file scan and indexes
- Try to evaluate predicates early

1 Left Deep Join Trees

- In **left-deep join trees**, the right-hand-side input for each join is a relation, not the result of an intermediate join.



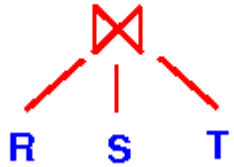
(a) Left-deep join tree



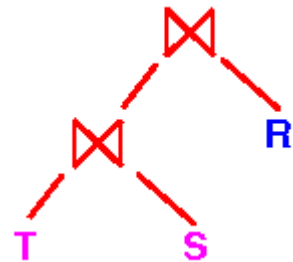
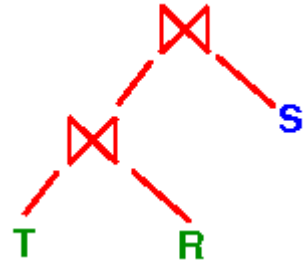
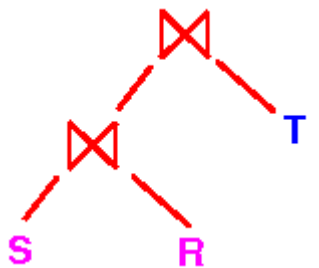
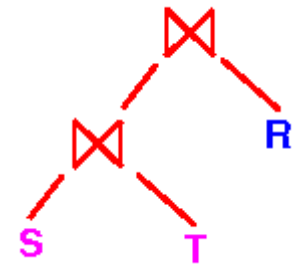
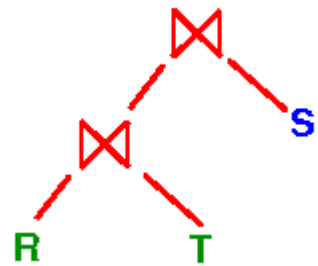
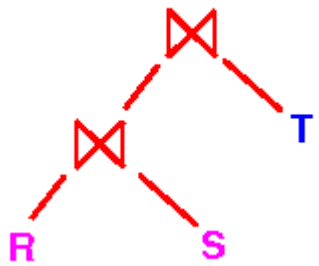
(b) Non-left-deep join tree

1 Left Deep Join Trees

Example: 3-way join:



Possible left-deep join trees



Number of *left-deep* join trees with N relations = $N !$ (factorial))

1 Why left Deep Join Trees?

Query plans based on:

- **Left-deep** join trees and
- **Commonly used join implementations (algorithms)**

tend to be more efficient because:

Left-deep join trees **interact very well** with **commonly used join (implementation) algorithms**

2 Plan Enumeration Algorithm

Idea: use dynamic programming

- For each subset of $\{R_1, \dots, R_n\}$, compute the best plan for that subset
- In increasing order of set cardinality:
 - Step 1: for $\{R_1\}, \{R_2\}, \dots, \{R_n\}$
 - Step 2: for $\{R_1, R_2\}, \{R_1, R_3\}, \dots, \{R_{n-1}, R_n\}$
 - ...
 - Step n: for $\{R_1, \dots, R_n\}$
- It is a bottom-up strategy
- A subset of $\{R_1, \dots, R_n\}$ is also called a *subquery*

2 Cost of Optimization

- **n=10 joins**
- **Finding the best join-order :** $(2(n-1))!/(n-1)! = 17,643,225,600$

With dynamic programming:

- Time complexity of optimization with bushy trees is **$O(3^n)$** = 59000
 - Space complexity is $O(2^n)$
-
- If only left-deep trees are considered, time complexity of finding best join order is **$O(n 2^n)$** = 10,240
 - Space complexity remains at $O(2^n)$
 - Cost-based optimization is expensive, but worthwhile for queries on large datasets (typical queries have small n, generally < 10)

3 Interesting Sort Orders

- Consider the expression $(r_1 \bowtie r_2) \bowtie r_3$ (with A as common attribute)
- An **interesting sort order** is a particular sort order of tuples that could be useful for a later operation
 - Using merge-join to compute $r_1 \bowtie r_2$ may be costlier than hash join but generates result sorted on A
 - Which in turn may make merge-join with r_3 cheaper, which may reduce cost of join with r_3 and minimizing overall cost
 - Sort order may also be useful for order by and for grouping

Interesting Sort Orders

- Not sufficient to find the best join order for each subset of the set of n given relations
 - must find the best join order for each subset, **for each interesting sort order**
 - Simple extension of earlier dynamic programming algorithms
 - Usually, number of interesting orders is quite small and doesn't affect time/space complexity significantly

Dynamic Programming Algo

- **Step 1: Enumerate all single-relation plans**
 - Consider selections on attributes of relation
 - Consider all possible access paths
 - Consider attributes that are not needed
 - Compute cost for each plan
 - Keep cheapest plan per “interesting” output order

Dynamic Programming Algo

-
- **Step 2: Generate all two-relation plans**
 - For each each single-relation plan from step 1
 - Consider that plan as outer relation
 - Consider every other relation as inner relation
 - Compute cost for each plan
 - Keep cheapest plan per “interesting” output order

Dynamic Programming Algo

-
- **Step 3: Generate all three-relation plans**
 - For each each two-relation plan from step 2
 - Consider that plan as outer relation
 - Consider every other relation as inner relation
 - Compute cost for each plan
 - Keep cheapest plan per “interesting” output order
 - **Steps 4 through n: repeat until plan contains all the relations in the query**

Commercial Query Optimizers

DB2, Informix, Microsoft SQL Server, Oracle 8

- Inspired by System R
 - Left-deep plans and dynamic programming
 - Cost-based optimization (CPU and IO)
- Go beyond System R style of optimization
 - Also consider right-deep and bushy plans (e.g., Oracle and DB2)
 - Variety of additional strategies for generating plans (e.g., DB2 and SQL Server)

Other Query Optimizers

Randomized plan generation

- Genetic algorithm
- PostgreSQL uses it for queries with many joins

Rule-based

- **Extensible** collection of rules
- Rule = Algebraic law with a direction
- Algorithm for firing these rules
 - Generate many alternative plans, in some order
 - Prune by cost
- Startburst (later DB2) and Volcano (later SQL Server)

Cost Based Optimization with Equivalence Rules

- **Physical equivalence rules** allow logical query plan to be converted to physical query plan specifying what algorithms are used for each operation.
- Efficient optimizer based on equivalent rules depends on
 - A space efficient representation of expressions which avoids making multiple copies of subexpressions
 - Efficient techniques for detecting duplicate derivations of expressions
 - A form of dynamic programming based on **memoization**, which stores the best plan for a subexpression the first time it is optimized, and reuses it on repeated optimization calls on same subexpression
 - Cost-based pruning techniques that avoid generating all plans
- Pioneered by the Volcano project and implemented in the SQL Server optimizer

Structure of Query Optimizers

- Many optimizers consider only left-deep join orders.
 - Plus heuristics to push selections and projections down the query tree
 - Reduces optimization complexity and generates plans amenable to pipelined evaluation.
- Heuristic optimization used in some versions of Oracle:
 - Repeatedly pick “best” relation to join next
 - Starting from each of n starting points. Pick best among these
- Intricacies of SQL complicate query optimization
 - E.g. nested subqueries

Structure of Query Optimizers

- Some query optimizers integrate heuristic selection and the generation of alternative access plans.
 - Frequently used approach
 - heuristic rewriting of nested block structure and aggregation
 - followed by cost-based join-order optimization for each block
 - Some optimizers (e.g. SQL Server) apply transformations to entire query and do not depend on block structure
 - **Optimization cost budget** to stop optimization early (if cost of plan is less than cost of optimization)
 - **Plan caching** to reuse previously computed plan if query is resubmitted
 - Even with different constants in query
- Even with the use of heuristics, cost-based query optimization imposes a substantial overhead.
 - But is worth it for expensive queries
 - Optimizers often use simple heuristics for very cheap queries, and perform exhaustive enumeration for more expensive queries

Optimizing Nested Subqueries

- Nested query example:

```
select name  
from instructor  
where exists (select *  
               from teaches  
               where instructor.ID = teaches.ID and teaches.year = 2007)
```

SQL conceptually treats nested subqueries in the where clause as functions that take parameters and return a single value or set of values

Parameters are variables from outer level query that are used in the nested subquery; such variables are called **correlation variables**

Optimizing Nested Subqueries

- Nested query example:

```
select name  
from instructor  
where exists (select *  
               from teaches  
               where instructor.ID = teaches.ID and teaches.year = 2007)
```

Conceptually, a nested subquery is executed once for each tuple in the cross-product generated by the outer level **from** clause

Such evaluation is called **correlated evaluation**

Note: other conditions in where clause may be used to compute a join (instead of a cross-product) before executing the nested subquery

Optimizing Nested Subqueries

- Correlated evaluation may be quite inefficient since
 - a large number of calls may be made to the nested query
 - there may be unnecessary random I/O as a result
- SQL optimizers attempt to transform nested subqueries to joins where possible, enabling use of efficient join techniques

Optimizing Nested Subqueries

- Nested query example:

```
select name  
from instructor  
where exists (select *  
               from teaches  
               where instructor.ID = teaches.ID and teaches.year = 2007)
```

can be rewritten as

```
select name  
from instructor, teaches  
where instructor.ID = teaches.ID and teaches.year = 2007
```

Optimizing Nested Subqueries

- In general, it is not possible/straightforward to move the entire nested subquery from clause into the outer level query from clause
 - A temporary relation is created instead, and used in body of outer level query

Optimizing Nested Subqueries

In general, SQL queries of the form below can be rewritten as shown

- Rewrite: **select ...**
 from L_1
 where P_1 and exists (select *
 from L_2 where P_2)
- To: **create table t_1 as**
 select distinct V
 from L_2 where P_2^1

 select ...
 from L_1, t_1 where P_1 and P_2^2
 - P_2^1 contains predicates in P_2 that do not involve any correlation variables
 - P_2^2 reintroduces predicates involving correlation variables, with relations renamed appropriately
 - V contains all attributes used in predicates with correlation variables

Optimizing Nested Subqueries

- In our example, the original nested query would be transformed to
 create table t_1 as
 select distinct ID
 from $teaches$
 where $year = 2007$

 select $name$
 from $instructor, t_1$
 where $t_1.ID = instructor.ID$
- The process of replacing a nested query by a query with a join (possibly with a temporary relation) is called **decorrelation**.
- Decorrelation is more complicated when
 - the nested subquery uses aggregation, or
 - when the result of the nested subquery is used to test for equality, or
 - when the condition linking the nested subquery to the other query is **not exists**,
 - and so on.